

AI Usage Disclosure

Julian Calvo. AAI6850, Module 15 capstone. August 2026.

I used Claude (Anthropic) across most of this project: the Solidity, the tests, the ML pipeline, the documentation, the diagrams, the slides, the CI config, and the first draft of the security review.

I checked the output by running it. The contracts compile under solc 0.8.24, 195 tests pass, coverage and Slither numbers come from the tools themselves, and the ML metrics reproduce under seed 6850. Reading the code rather than accepting it turned up four bugs in earlier revisions: provider bonding that did nothing, an exclusivity gap, an emission accounting error, and a compliance bypass. Those are in `audit/security_audit_report.md`, which is my own review, not an independent audit.

Third-party code: OpenZeppelin Contracts 5.6.1, the Synthetix reward accounting pattern, Hardhat, ethers.js, scikit-learn, and the rest of `package.json` and `ml/requirements.txt`, all under their own licences. Published methods and market data are cited in the reports that use them.

The testnet deployment and the cross-chain run are real. Transaction links are in `docs/DEPLOYMENT.md`.